Same model, faster decoding

Hybrid + energy • identical checkpoint and two refinement rounds • 8,192-token prefix

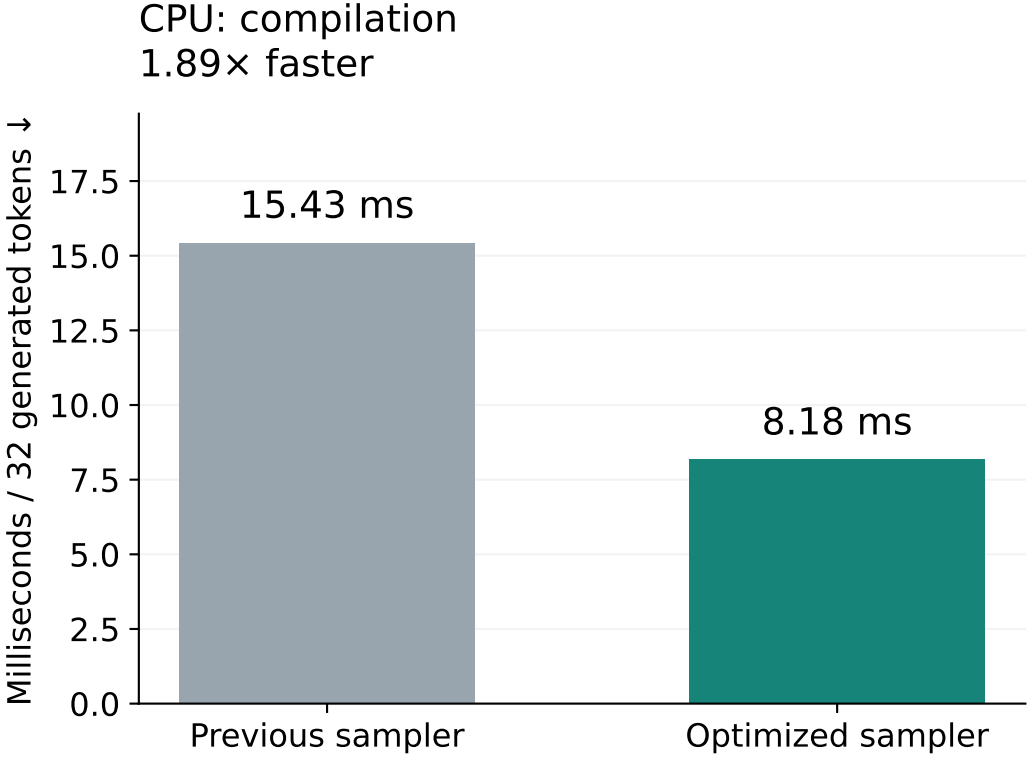

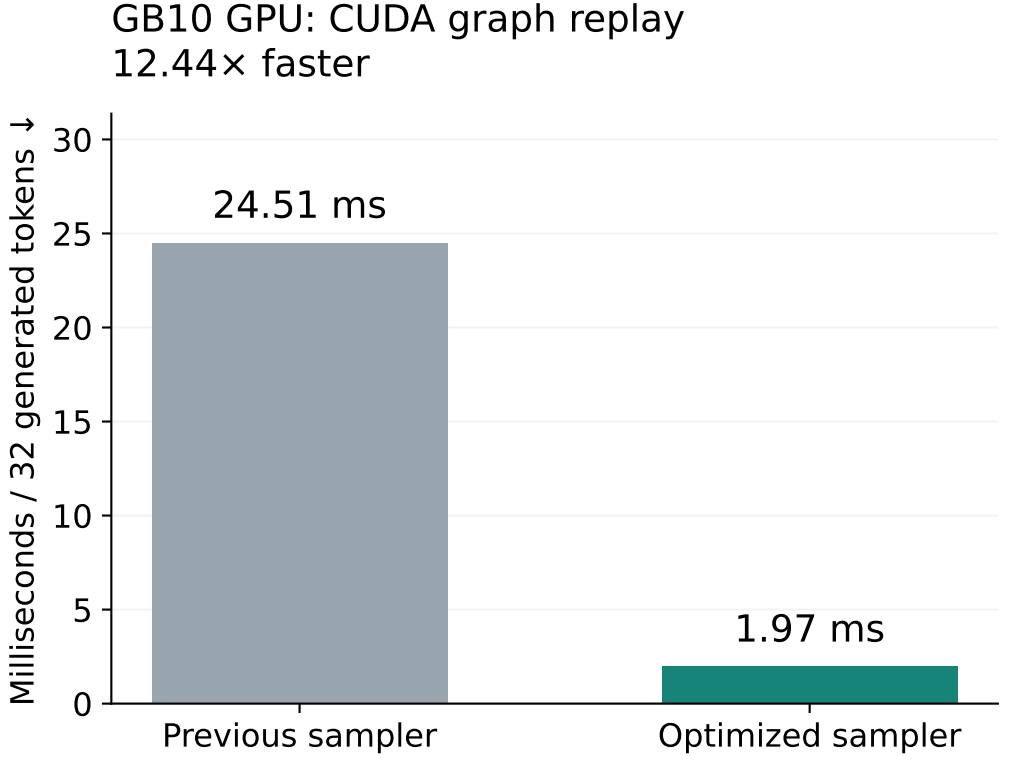

Medians of nine alternating runs; warmed generation including token return to CPU.
Prefill, compilation and graph capture excluded. Shared GPU (3% → 44% utilization); no training-speed claim.