Compact memory is the clearest result

~112K parameters • recall: three training seeds • timing: one checkpoint per model

Recall through distractions

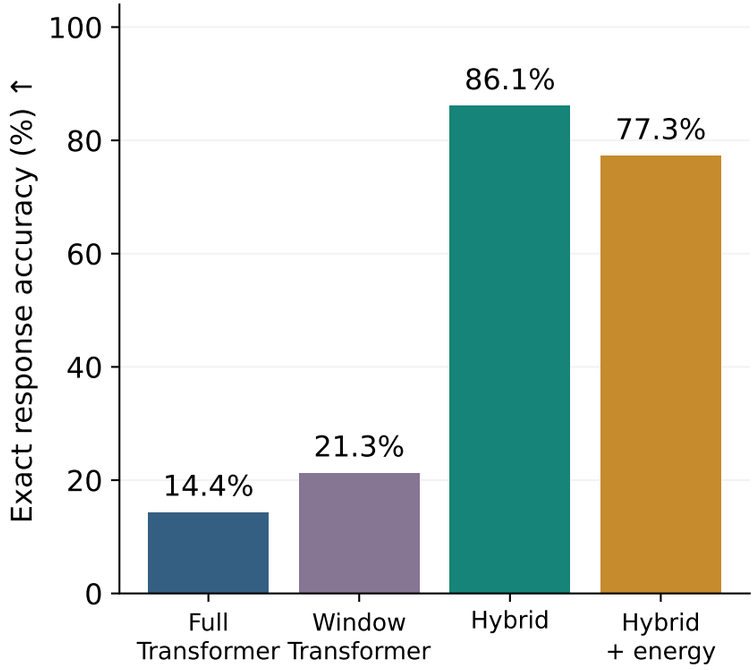

32–128 distractor facts; up to 520 bytes
72 test prompts × 3 seeds per model

Small carried state at 8K

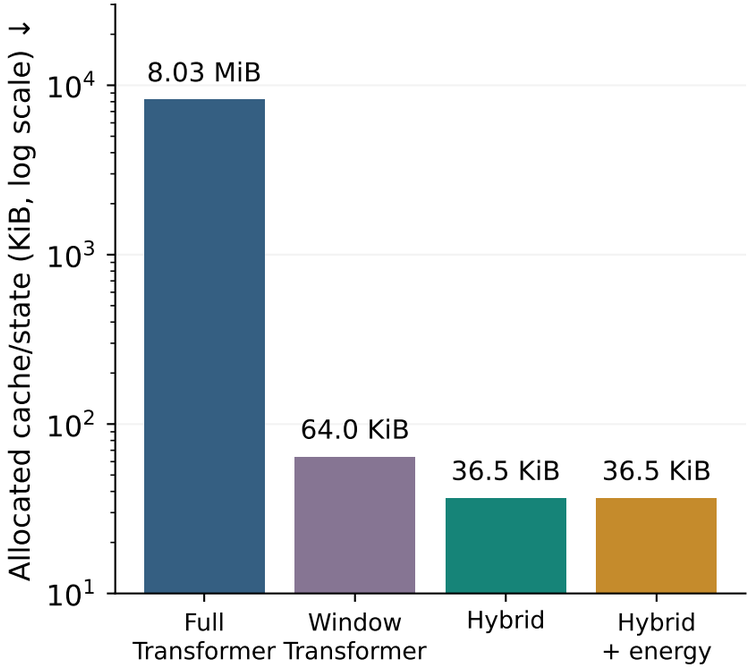

8,192-token prefix + 32 output capacity
Excludes weights and runtime workspaces

Faster long-prefix CPU decode

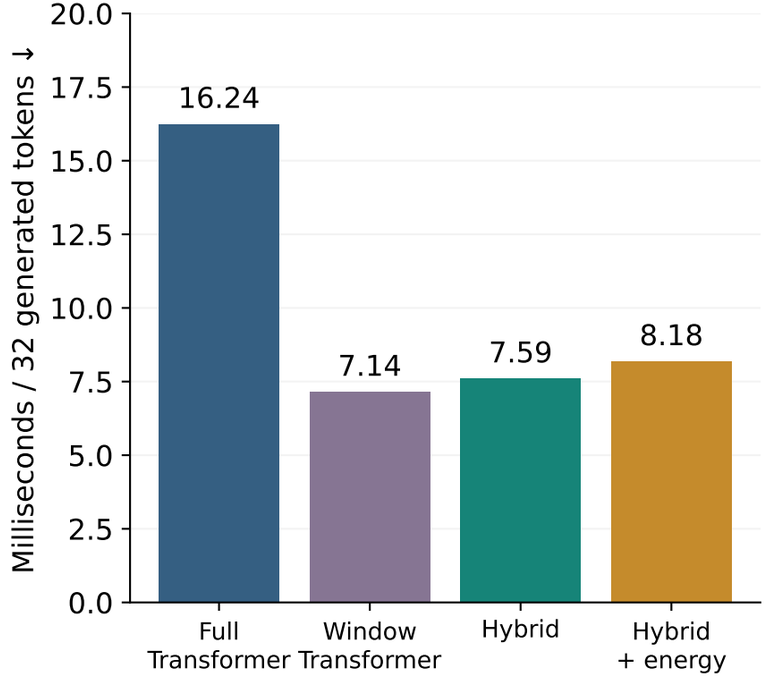

One CPU core; all arms compiled
Warmed decode only; nine repetitions

Energy refinement has not improved recall. The 8K measurements test execution cost, not 8K recall accuracy.