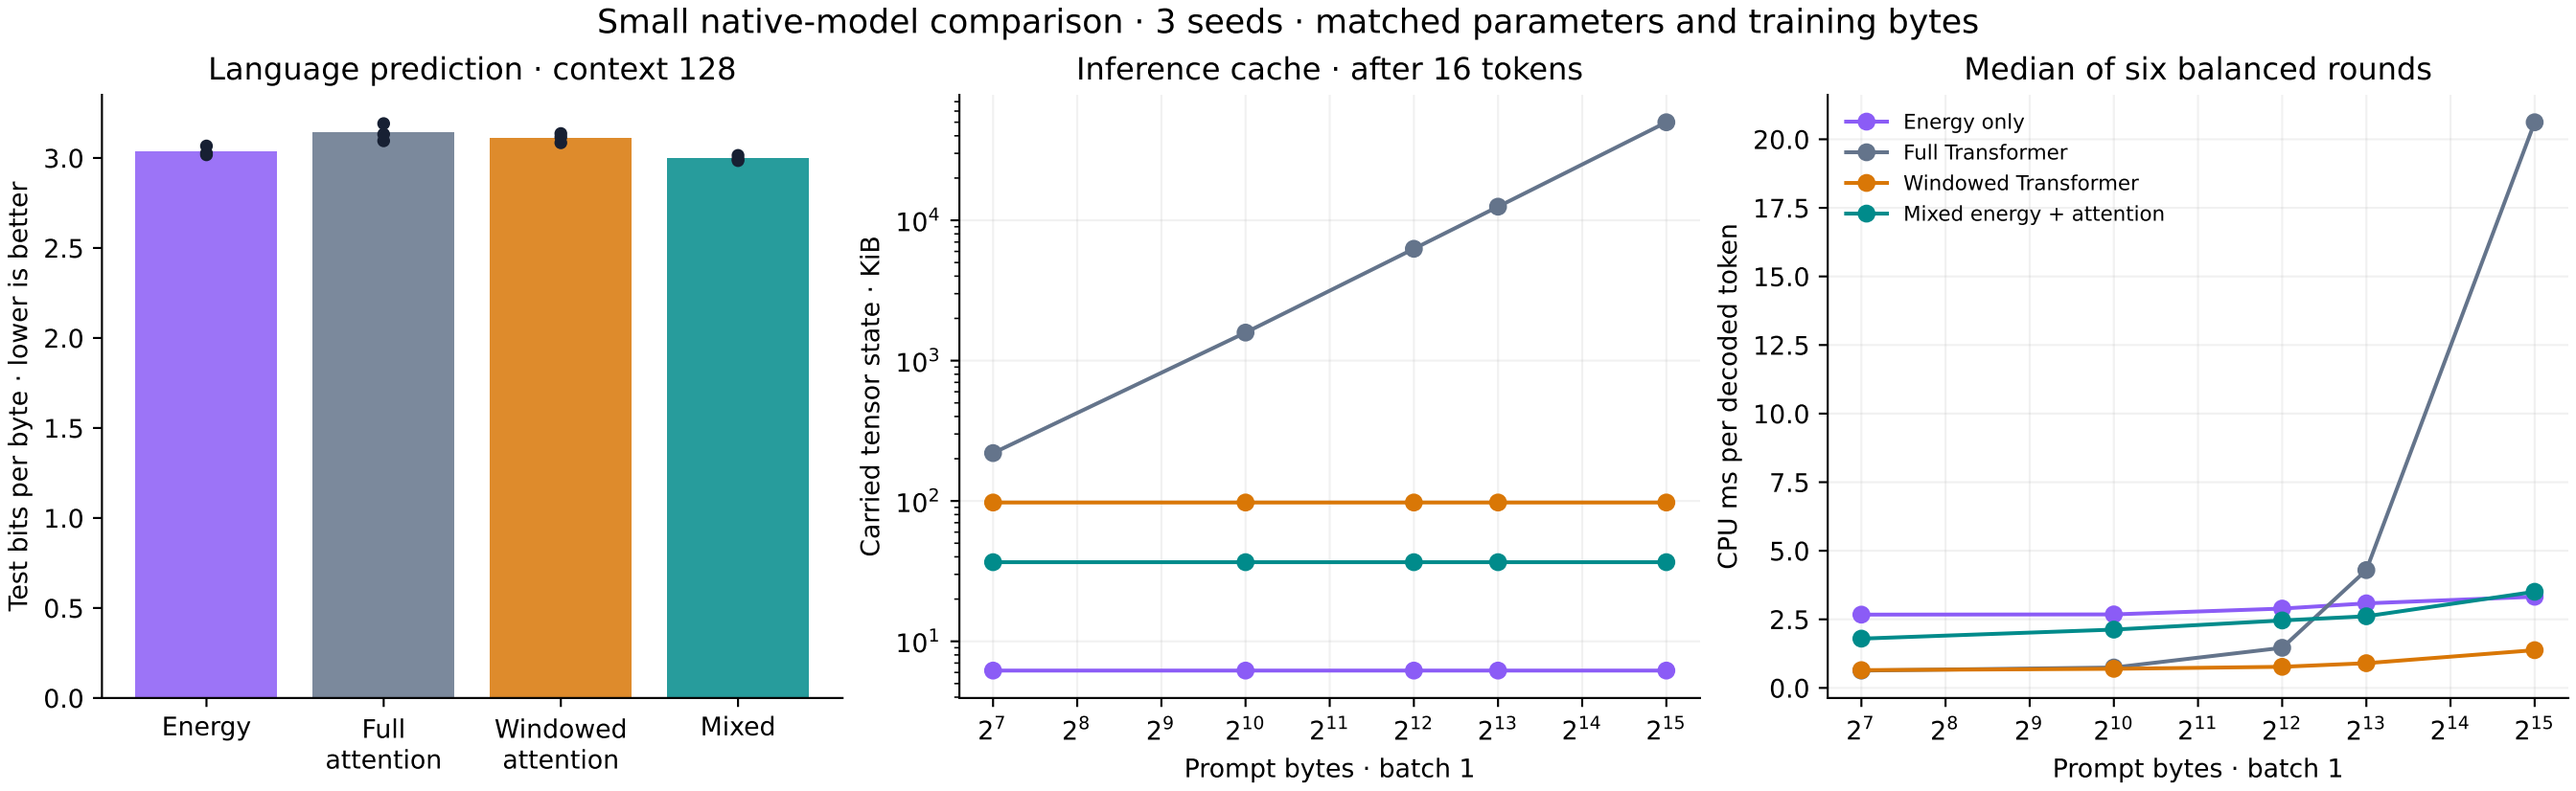
Small native-model comparison · 3 seeds · matched parameters and training bytes
Language prediction · context 128
Test bits per byte · lower is better
0.0
0.5
1.0
1.5
2.0
2.5
3.0
Energy
Full attention
Windowed attention
Mixed
Inference cache · after 16 tokens
Carried tensor state · KiB
10^1
10^2
10^3
10^4
2^7 2^8 2^9 2^{10} 2^{11} 2^{12} 2^{13} 2^{14} 2^{15}
Prompt bytes · batch 1
Median of six balanced rounds
CPU ms per decoded token
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0
Energy only
Full Transformer
Windowed Transformer
Mixed energy + attention
2^7 2^8 2^9 2^{10} 2^{11} 2^{12} 2^{13} 2^{14} 2^{15}
Prompt bytes · batch 1