Frozen REWT test results · October 10, 2026

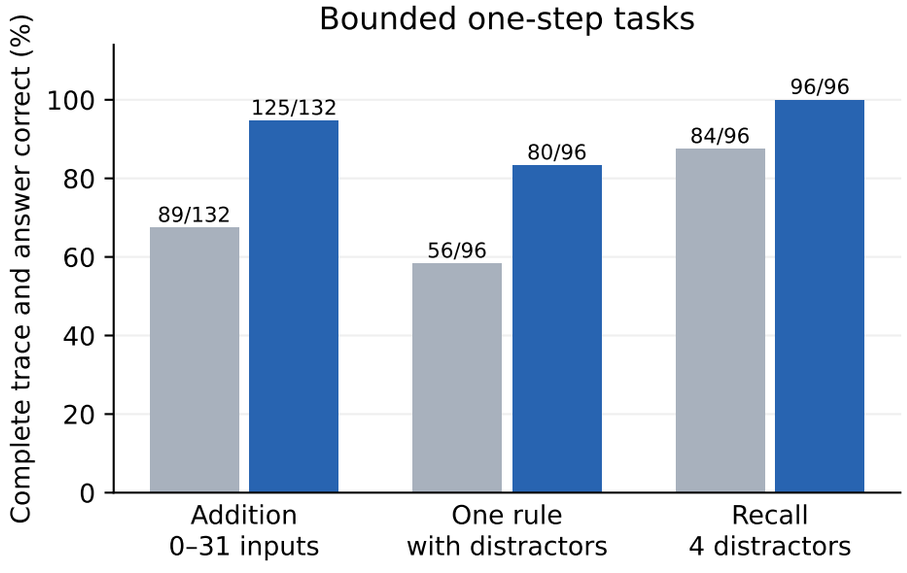

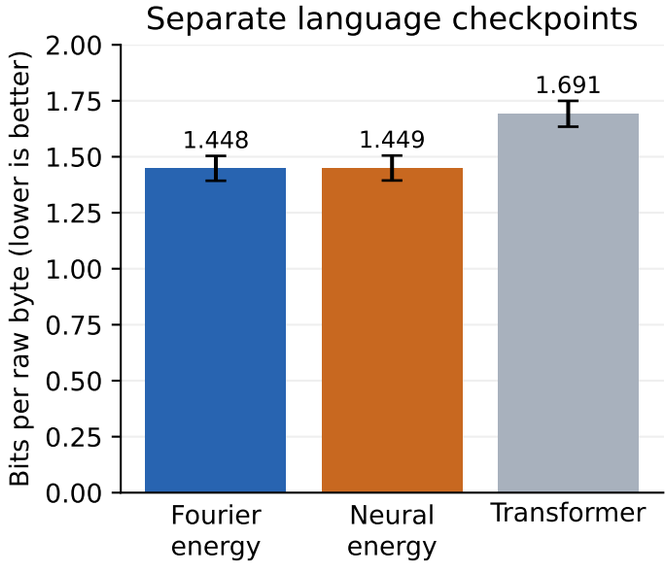

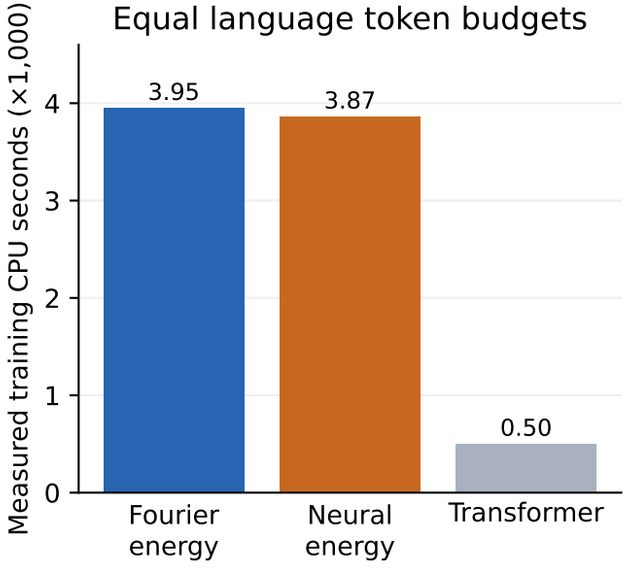

One initialization lineage per model; task and language models are separate. Arithmetic has 66 semantic pairs in two wordings.
Language bars: 64 unseen story windows; intervals resample documents and do not represent training-seed variability.
All 90 generated story continuations failed the exploratory usefulness rubric. Lower predictive loss does not establish coherent prose.