Streaming REWT: measured wins and remaining costs

~112K parameters · three training seeds · fresh held-out tests · independent timing

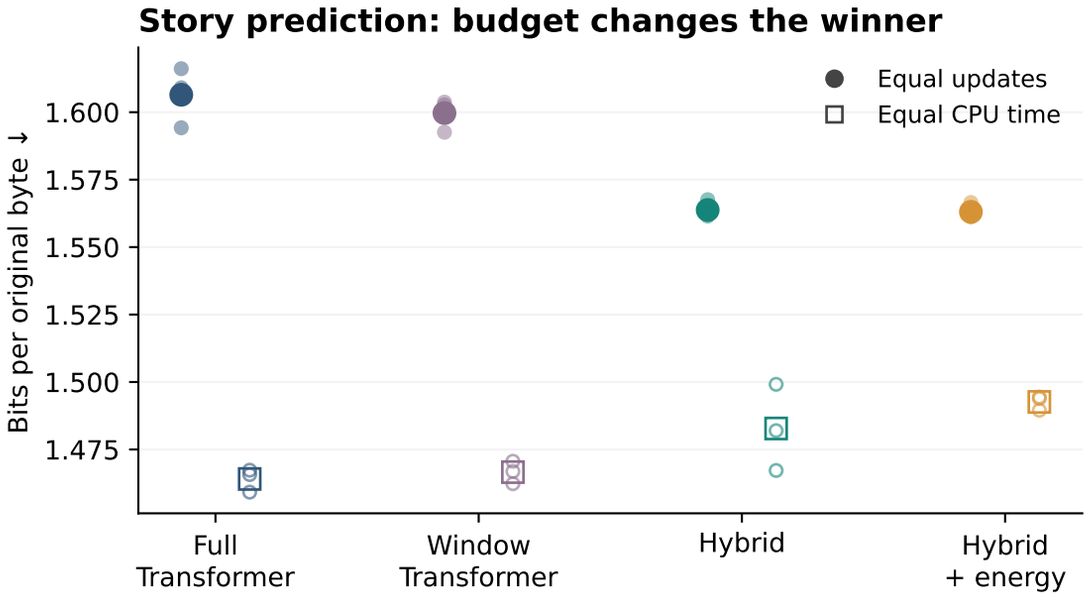

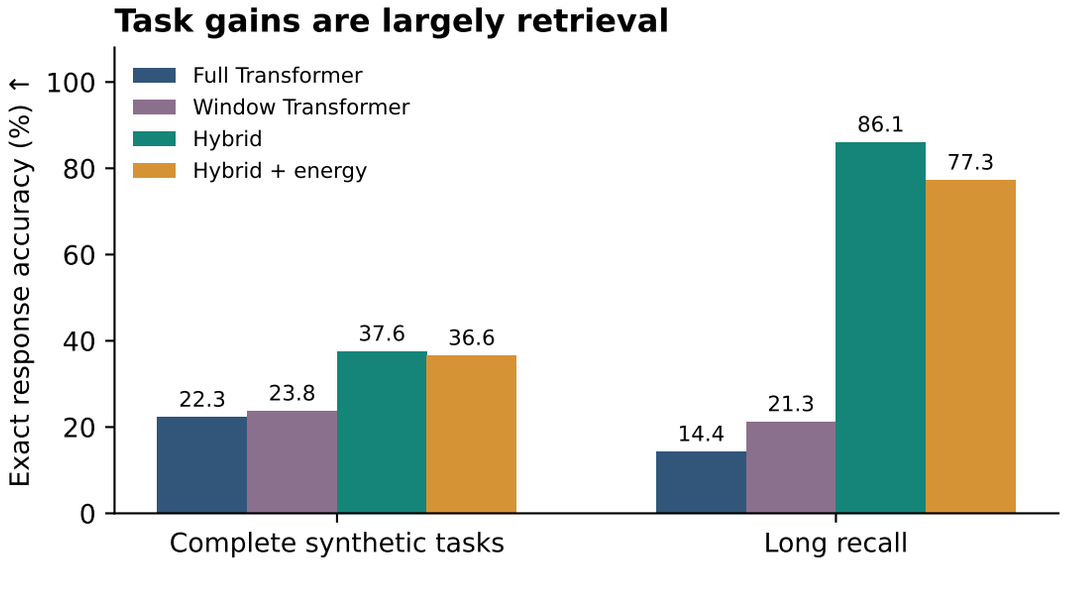

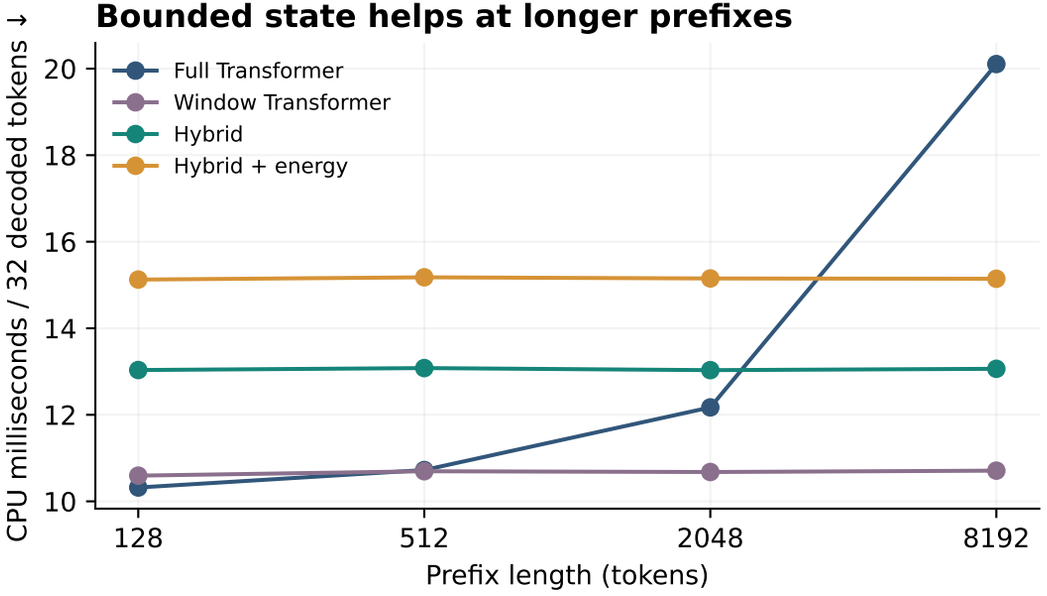

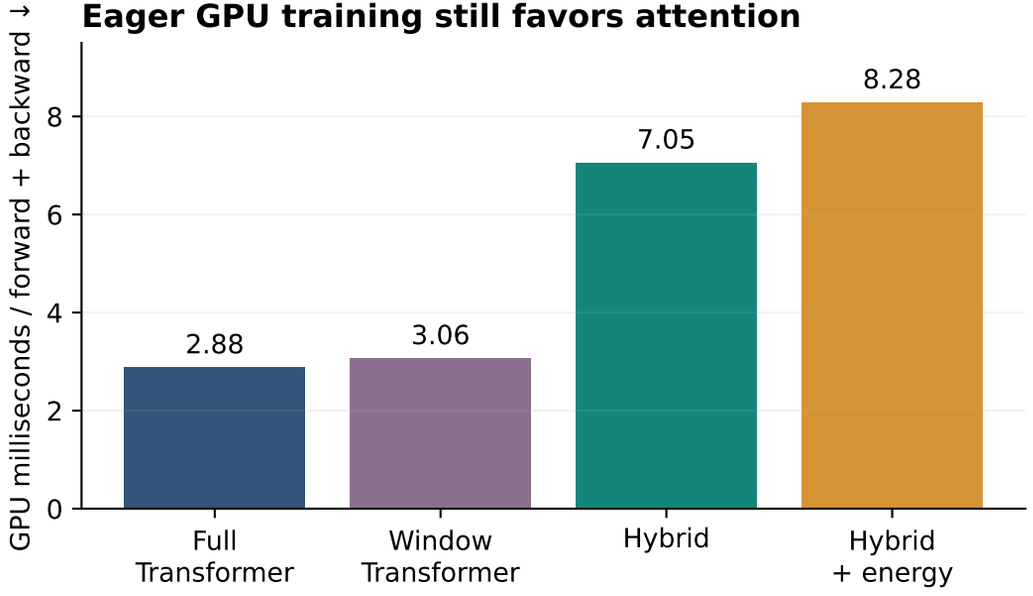

CPU: one pinned core, preallocated caches. GPU: shared GB10, float32, 24×128, eager execution. Small-model results; no general superiority claim.