Whole-word prediction on fresh TinyStories

1,024 unseen stories · 65,536 target positions · 3 training seeds · about 1.40M parameters

Matched parameters and training tokens

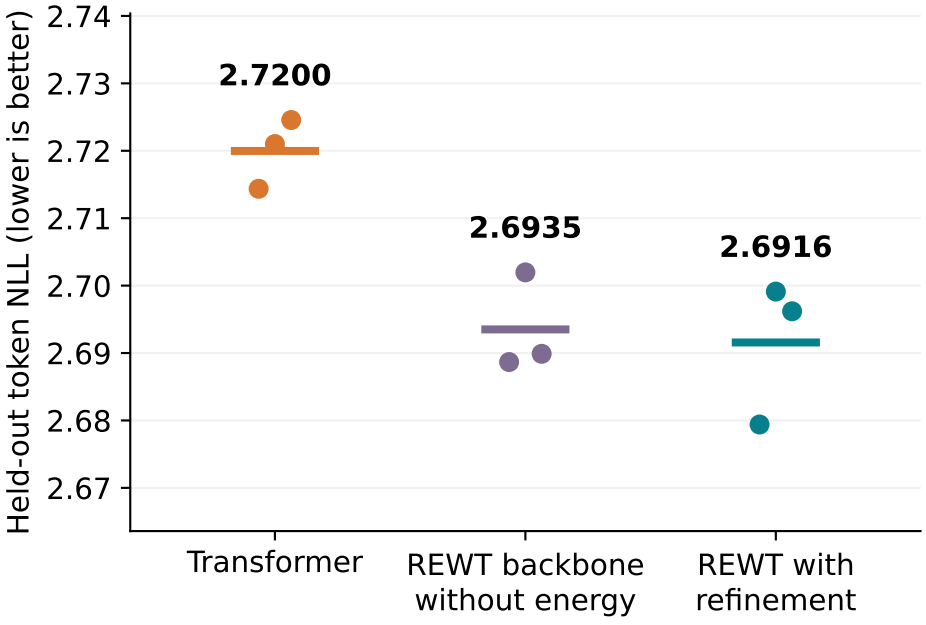

More refinement, same trained weights

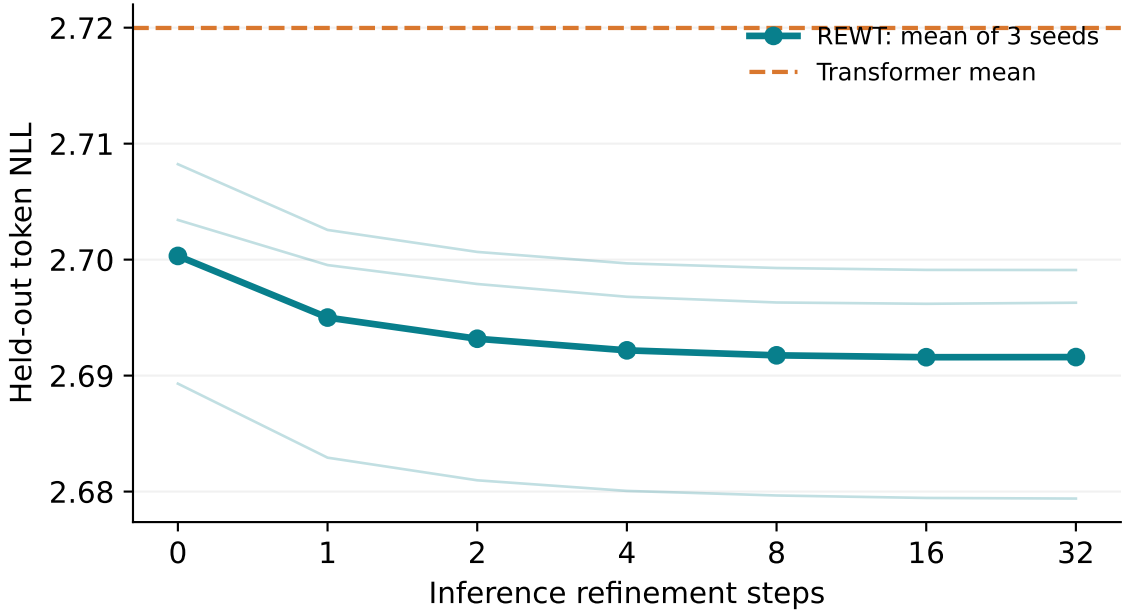

12.288M training-token exposures per fit. Checkpoint/depth selection uses validation.
NLL includes punctuation and special tokens; training FLOPs are not matched.